
Store Sales - Time Series Forecasting

Ritu Pravakar
pravakar@usc.edu

Charlene Yuen
yuenchar@usc.edu

Sreenidhi Iyengar Munimadugu
munimadu@usc.edu

Mahek Savani
msavani@usc.edu

Abstract

The paper outlines the process of developing a solution and the final outcome from the Kaggle Competition Store Sales. Given dates, store and product information, promotions, sales numbers, as well as additional information such as gas prices and holidays, we utilized what was given to make a prediction on the sales of product families at different Favorita stores in Ecuador. Below we detail our progress with different pre-processing methods, feature engineering, and models.

1 Introduction

In current society, there is a multitude of applications for forecasting, ranging from the weather to the economy. One particular application is forecasting store sales, allowing grocery stores to anticipate how much inventory to purchase. Given this information, grocers can stock exactly how much is in demand from customers without running out or having extra. Therefore, there lies great potential in reducing food waste, as well as saving time and other resources. For the purposes of the competition, we outline our thought process of feature engineering and developing a model for forecasting store sales.

2 Data Preprocessing

The provided data included dates, store and product information, promotions, sales numbers, as well as additional information such as gas prices and holidays. To begin, we read in the data from the provided CSVs and perform some initial cleaning. We extract the date from the index and create a new column 'days_from_2013' for easy shifting of the time series data. We also group the data by store and product family to calculate the mean sales for each group. In the data containing daily oil prices, there were some NaN values, which we filled with interpolation backward based on the next valid value, which manages to capture the overall trend yet fill in missing points. We later also drop NaN values to keep the most relevant data. The same happened with the data containing transactions from different dates and stores, so a similar filling of missing values was done by setting null values to 0. We also grouped transactions by their store number and date which made taking moving averages over transactions and merging with the rest of the dataframe easier.

3 Feature Engineering

Next, we perform feature engineering on the data to create additional features that may help improve the accuracy of our forecasts.

3.1 Oil

The Ecuadorian economy has been shown to be dependent on exports, one of which is oil, making up a hefty 36% of total exports [1]. Due to this dependency, it follows that when the economy is poor, high oil prices would result in plummeting sales. Correspondingly, when plotting oil prices and sales with respect to each other, we found a negative trend (Figure 1).




Figure 1: Correlation between oil and sales




Figure 2: 100 oil lags and partial autocorrelation

We then took the moving average over a week, two weeks, and a month to capture any possible correlations to sales. We also used the partial autocorrelation plot of oil lags from 1 to 100 and found that there were many lag features that contributed to new correlations, so we simply picked lags with the largest values (Figure 2).

3.2 Holidays and events

We expected holidays to be correlated with an increase in sales, particularly near the end of the year (Figure 3).



Figure 3: Total sales by year

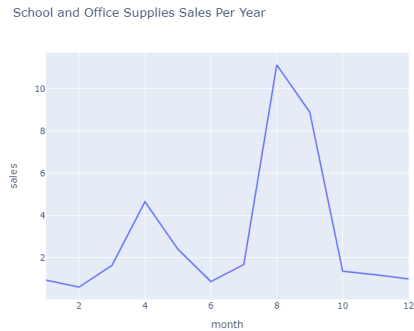


Figure 4: Sales of school by month

Thus, we separated workdays from holidays based on if the day is of type Bridge, Transfer, Additional and untransferred, as well as a holiday that is untransferred. One-hot encoding was also done for events, and features were tested for certain occurrences such as the New Year, Black Friday, the earthquake, or Easter. However, overall they did not seem to make a large improvement so not all of these features were included in our final solution. Other patterns found included sales being higher on weekends and sales of school supplies spiking during the back-to-school seasons of April-May and August-September, which is shown in Figure 4.

These findings led us to add features dictating if the date was a school season or the weekend which both correlate with higher sales.

3.3 Stores

From the given store data, we identified features such as certain stores being the only store of that type or a store that opened later, which could identify patterns in the data (Figure 5).




Figure 5: Count of different stores

57

58 Merged with training data, we can also categorize stores and compare the average sales of each city,
 59 state, and cluster. Our findings concluded that the top cities with the best overall sales were Quito,
 60 Cayambe, Ambato, and Dalue; the states with the best sales were Pichincha, Tungurahua, Loja, and
 61 and El Oro; the clusters with the best sales were 5, 14, 8, and 11. Plots can be found below in Figure 6.
 62




Figure 6: Total sales by city, state, cluster from 2013-2017

63 With this knowledge, we tested new features for the best sales from each category, which gave a
 64 small amount of improvement though it did not make it into our final solution due to time constraints.
 65 Overall, the stores data gives good information on patterns based on location and groupings of the
 66 stores.

67 3.4 Transactions

68 From the transaction data, intuitively transactions should be correlated with sales, which we found
 69 had a correlation of 0.8. Apart from that, we filled in missing data with zeroes and grouped the data
 70 by store number and date so we could take moving averages and lags over the transactions to capture
 71 any similarities to sales. We also identified the total number of items on promotion in store, which
 72 generally proportional to sales excluding an outlier as shown in Figure 7.

73 Additional one-hot encoding was done for other time-related features such as days of the week, days,
 74 months, years, etc. For example, weeks can hold patterns in average transactions (Figure 8) and
 75 seasonal patterns may be caught by these features, such as sales rising towards the end of the year
 76 (Figure 3).

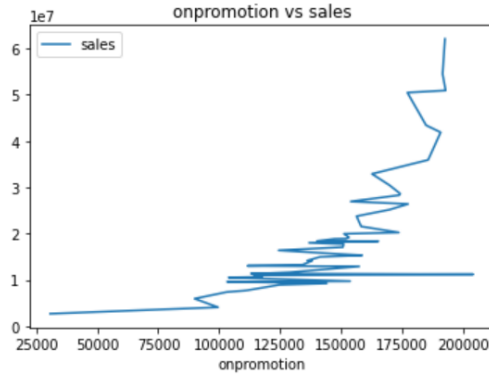


Figure 7: Items on promotion vs sales




Figure 8: Average transactions per week

77 4 Models

78 4.1 XGBoost

79 In XGBoost a common decision is made by all decision trees where the results of the previous tree
80 affect the results of the next decision tree. It is one of the fastest decision tree algorithms and predicts
81 more accurately than a single decision tree.

82 4.2 Random Forest

83 Decision trees select important features and can give the relationship between features and prediction.
84 This helps in a better interpretation of the model. A random forest is an ensemble of trees where
85 at each node, the best split is chosen among the randomly selected subset of features. For the sales
86 forecast, the previous value is used by the predictor variable.

87 4.3 Hybrid Models

88 We learned the components of the model in an iterative way using a hybrid model. Regression
89 models like linear and ridge can extrapolate trends while ensemble methods like random forests
90 cannot. Random forests can learn the interactions between different features. For our model, we first
91 used ridge regression to extrapolate. In the next step, we transformed the target and removed the
92 trend. We applied random forest to the detrended residuals, which was the difference between the
93 target the model was trained on and the predictions it made [3].
94

Algorithm 1	Algorithm 2	Mean squared log error	Root mean squared log error
ExtraTreesRegression	RandomForest	0.00000289	0.00032
RandomForest	Linear	8.54	0.52
RandomForest	Ridge	8.48	0.51
Linear	RandomForest	6.80	0.47
Ridge	RandomForest	6.86	0.47

Table 1: Comparison of Mean squared log error and Root mean squared log error

95

96 4.4 LightGBM

97 LightGBM [4] is an ensemble model, where the result of all trees are combined in the output.
98 XGBoost [5] pre-sorts the data by features and then traverses the data set to find the best segmentation

99 point on features. This increases memory usage and time, hence LightGBM model is optimized in
100 Gradient Facilitated decision tree algorithm.

101 We also compared the performance of the model using XGBoost, LightGBM, Linear Regression, and
102 Random Forest.

103 **5 Final Results**

104 After modeling hybrid models we found out that it will take a lot of hyperparameter tuning for all the
105 models at the same time to lower the RMSLE loss, so instead we decided to tune the hyperparameters
106 of the Random Forest Regressor.

107 We tune the hyperparameters with Optuna, a hyperparameter tuning library that is used to find the
108 best hyperparameters for the model. We split the data into training and validation sets, and apply log
109 transformation to the sales data. We then fit the model on the training data, and make predictions on
110 the validation set.

111 The parameters we tried to tune for random forest regressor are as follows:

112 Criterion: The function to measure the quality of a split. We used RMSE for this parameter

113 Bootstrap: Whether bootstrap samples are used when building trees. If False, the whole dataset is
114 used to build each tree.

115 Max Depth: The maximum depth of the tree. If None, then nodes are expanded until all leaves are
116 pure or until all leaves contain less than min_samples_split samples.

117 Max Features: The number of features to consider when looking for the best split.

118 Max Leaf Nodes: Grow trees with max_leaf_nodes in the best-first fashion. Best nodes are defined
119 as relative reduction in impurity. If None then an unlimited number of leaf nodes.

120 N Estimators: The number of trees in the forest.

121 Min Samples Split: The minimum number of samples required to split an internal node.

122 Min Samples Leaf: The minimum number of samples required to be at a leaf node.

```
123 criterion: 'squared_error',  
124 bootstrap: [True, False],  
125 max_depth: 1, 10000,  
126 max_features: ['auto', 'sqrt', 'log2'],  
127 max_leaf_nodes: 1, 10000,  
128 n_estimators: 30, 5000,  
129 min_samples_split: 2, 100,  
130 min_samples_leaf: 1, 50
```

131 The best parameters after the tuning are as follows:

```
132 'criterion': 'squared_error',  
133 'bootstrap': 'False',  
134 'max_depth': 9733,  
135 'max_features': 'auto',  
136 'max_leaf_nodes': 4730,  
137 'n_estimators': 700,  
138 'min_samples_split': 3,  
139 'min_samples_leaf': 8
```

140 We found out from parameter tuning that increase in the number of trees (n_estimators) in the Random
141 Forest Regressor, gave the highest improvement in reducing the RMLSE error/loss.

142 Also, we found an interesting observation that after a certain point adding the number of trees
143 didn't reduce the RMLSE error. We found out that the decrease of RMLSE Loss plateau at 700-800
144 estimators beyond that increase just overfits the model without any substantial decrease in RMLSE.
145 Hence we choose 700 n_estimators to be a good choice for the data.

146 After training the model and making predictions on the validation set, we evaluate the performance
147 using the mean squared logarithmic error (MSLE). We then apply the model to the test data, and make
148 final predictions. We post-process the predictions by applying the inverse of the log transformation,
149 and save the results to a CSV file.

150 The below figure gives the idea of an increase in the number of trees vs a decrease in RMLSE loss
151 (Figure 9).

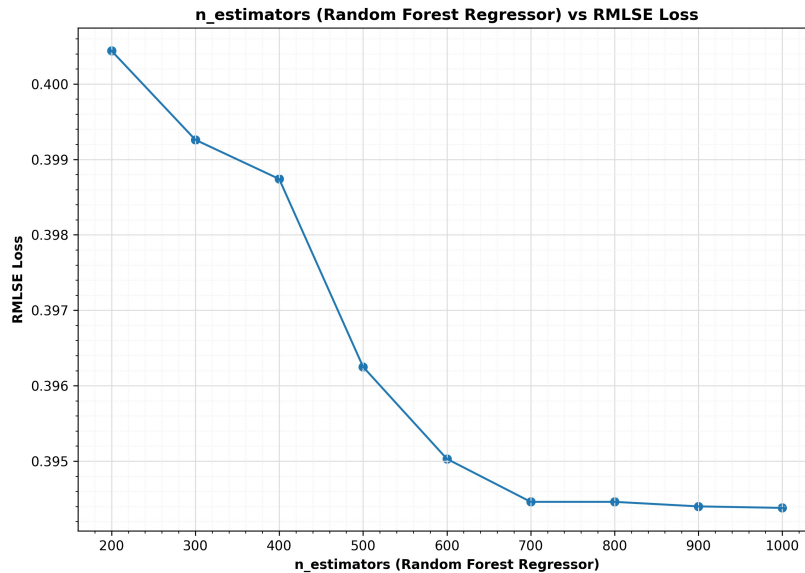


Figure 9: Hyperparameter tuning results

152 6 Running instructions

153 To run the code and generate the results, first ensure that all necessary libraries are installed and
154 imported. Upload as a Kaggle notebook and run the code cells in the order they are presented in the
155 notebook. The final predictions will be saved to the 'submission.csv' file.

156 References

- 157 [1] Ecuador's economy and oil: <https://www.bmz.de/en/countries/ecuador/economic-situation-51990>
- 158 [2] Kaggle Time Series Tutorial: <https://www.kaggle.com/learn/time-series>
- 159 [3] <https://www.kaggle.com/code/ryanholbrook/hybrid-models>
- 160 [4] Ke, Guolin, Qi Meng, Thomas Finley, Taifeng Wang, Wei Chen, Weidong Ma, Qiwei Ye, and Tie-Yan Liu.
161 "Lightgbm: A highly efficient gradient boosting decision tree." Advances in neural information processing
162 systems 30 (2017).
- 163 [5] Chen, Tianqi, and Carlos Guestrin. "Xgboost: A scalable tree boosting system." In Proceedings of the 22nd
164 acm sigkdd international conference on knowledge discovery and data mining, pp. 785-794. 2016.